A Simple Guide to Stochastic Processes

M.S. Corey Zhang

University of Delaware
West Chester University of Pennsylvania

April 2024

An introductory exploration of stochastic processes, covering random walks, Markov chains, martingales, Brownian motion, and Ito calculus, with applications in data science and computational mathematics.

Abstract

Stochastic processes help us understand systems that change over time in unpredictable ways. This article explains the basics of these processes using clear examples and real-world connections. We'll cover simple random walks, Markov chains, martingales, and Brownian motion, then introduce Ito calculus, a tool used in fields like finance and physics to handle random changes.

1 What Are Stochastic Processes?

A stochastic process is like a sequence of random events tracked over time. Imagine flipping a coin every minute and recording the results. That is a simple example. These processes can happen in steps, such as $X_0, X_1, X_2, \ldots$, or continuously, like a stock price moving all day, $t \geq 0$. Think of it as a way to predict the paths a system might take, even there is randomness involved.

These ideas are used everywhere from stock prices to populations grow. They let us make good guesses about what might happen next.

2 Random Walks: A First Look

A random walk is one of the easiest stochastic processes to understand. Picture someone standing at zero on a number line, taking a step left or right with equal chance, like flipping a coin. We call this a one-dimensional symmetric random walk. Mathematically, the position X_t is the sum of steps $Y_i \in \{1, -1\}$, so $X_t = \sum_{i=1}^{t} Y_i$, with $X_0 = 0$. On average, the walker stays near zero, but over time, their position can vary, grows with t, which could make the walk's behavior hard to analyze because it gets wilder as time goes on. To keep things manageable, we divide by $\sqrt{t}$, creating the normalized process $X_t/\sqrt{t}$. The central limit theorem tells us this scaled process starts to look like a bell-shaped curve, known as a standard normal distribution, with variance of 1.

Random walks are great for studying things like waiting lines or stock market risks. For example, what's the chance the walker reaches $+100$ before hitting -50? The answer is $50/(100+50) = 1/3$. This kind of question helps us understand the chance of losing money in investments.

3 Markov Chains: Predicting the Next Step

Markov chains have the systems where the next step only depends on where you are now, not how you got there. For example, if you're at a certain point, the chance of moving to another point is always the same, regardless of the path. We can call this memory-less, and we can write this as:

$$P(X_{t+1} = s \mid X_0, \ldots, X_t) = P(X_{t+1} = s \mid X_t)$$

We use a transition matrix P to show the chances of moving between states. Each entry $P(i, j)$ is the probability of going from state i to state j. Over time, we can use this matrix to predict future steps, such as we can calculate P^n for n steps.

Many Markov chains usually converge into a steady matrix, called a stationary distribution π, where $\pi P = \pi$. This tells us how the system become in a long run, no matter where it started. It's like predicting weather patterns over years based on current conditions.

4 Martingales: Fair Games

Martingales model systems is like a fair game that you're not expected to win or lose value from one step to the next. It means a process X_t is a martingale if the expected value of the next step equals the current. We can write this as:

$$\mathbb{E}[X_{t+1} \mid X_0, \ldots, X_t] = X_t$$

Unlike Markov chains, which depend only on the current state, martingales consider all past states to compute this expectation. This fairness means you can't expect to come out ahead by choosing when to stop. A rule called the optional stopping theorem says under certain conditions, if you stop a martingale at a specific time τ, the expected value stays the same: $\mathbb{E}[X_\tau] = \mathbb{E}[X_0]$. This shows you can't beat the system just by picking the right moment to quit. These help us understand things like gambling or investment strategies.

5 Brownian Motion: Continuous Randomness

Brownian motion, also called the Wiener process. It is like a random walk but in continuous time. Unlike a random walk, which takes discrete steps (+1 or -1) at fixed times, Brownian motion moves smoothly over all times $t \geq 0$ with changes that follow a normal distribution. For times $s < t$, the change $B(t) - B(s) \sim \mathcal{N}(0, t-s)$ means the movement depends on how much time passes. Its paths are continuous but nowhere differentiable. Brownian motion is the scaling limit of a random walk: if you take a random walk with smaller and smaller steps (size δx) over shorter time intervals (δt), scaling so that $\delta x = \sqrt{\delta t}$, the process converges to Brownian motion as $\delta t \to 0$.

This process is important in modeling stock prices in finance or particles bouncing in physics. One useful feature is its quadratic variation. Imagine breaking time into tiny intervals, like splitting a second into many small pieces, and measuring how much Brownian motion changes in each piece. If you square these tiny changes and add them up, the total equals the elapsed time, the duration of the period you're observing, written as $dB_t^2 = dt$. For example, if you watch Brownian motion for one second, the sum of these squared changes equals 1, matching the time duration. This property shows how unpredictable Brownian motion is and is crucial for Ito calculus, which rely on this to handle this motion.

6 Maximum Values and Paths

Brownian motion has interesting patterns. Let $M(t) = \max_{s \leq t} B(s)$ represent the highest value Brownian motion reaches from time 0 to time t. For example, if we model a stock price, $M(t)$ is the peak price during that period. The chance that this maximum goes above a positive level $a > 0$ is twice the chance that the process's value at time t is above a:

$$P(M(t) > a) = 2P(B(t) > a)$$

This happens because Brownian motion is equally likely to move up or down. If it hits a at some point, it might stay above or fall below by time t, and this symmetry doubles the probability. Imagine a stock price that has chance hits 50 during a day is twice the chance it's above 50 at the day's end. This is useful for things like financial options, where we care about whether a price crosses a certain level.

7 Ito Calculus: Handling Random Changes

Regular calculus doesn't work for Brownian motion because its paths are nowhere differentiable, meaning they are continuous but not have a well-defined slope. Ito calculus is a special tool that helps. It includes Ito's Lemma, which is like a chain rule for random processes. If we apply a smooth function f, like $f(x) = x^2$, to the Brownian motion B_t, Ito's Lemma tells us how $f(B_t)$ changes over time:

$$df(B_t) = f'(B_t)\,dB_t + \frac{1}{2}f''(B_t)\,dt$$

f must be smooth (with continuous first and second derivatives), even though B_t is not. The extra term with f'' accounts for the randomness of Brownian motion's irregular paths.

For example, let's take $f(B_t) = B_t^2$. The first derivative is $f'(B_t) = 2B_t$, and the second derivative is $f''(B_t) = 2$. Applying Ito's Lemma:

$$d(B_t^2) = 2B_t dB_t + \frac{1}{2} \cdot 2 \cdot dt = 2B_t dB_t + dt$$

The change in B_t^2 has two components: $2B_t dB_t$, driven by the random movement of B_t and dt, which increases because the squared B_t sum to the time elapsed. If B_t represents stock price fluctuations, B_t^2 will be volatility, and the dt term reflects the steady accumulation over time.

Consider another function, $f(x) = e^x$, where $f(B_t) = e^{B_t}$ models exponential growth, such as a stock price in option pricing models. We get $f'(x) = e^x$, $f''(x) = e^x$, so:

$$d(e^{B_t}) = e^{B_t} dB_t + \frac{1}{2} e^{B_t} dt$$

The term $\frac{1}{2} e^{B_t} dt$ contributes to the growth of e^{B_t} because e^{B_t} is always positive and the second derivative $f''(x) = e^x$ is positive. Additionally, the exponential function's convex shape causes larger increases in e^{B_t} when B_t rises compared to smaller decreases when B_t falls by the same amount. For example, if B_t increases by 1, e^{B_t} multiplies by approximately 2.718, but if B_t decreases by 1, e^{B_t} divides by about 2.718, resulting in a smaller change. This asymmetry, combined with the positive $\frac{1}{2} e^{B_t} dt$ term from the quadratic variation causes e^{B_t} tend to increase over time. The choice of f depends on the application: x^2 for volatility, e^x for price growth, or other functions for physical properties such as energy or signal patterns. Ito's Lemma effectively handles the unpredictable nature of Brownian motion, making it a useful tool in finance for applications like the Black-Scholes formula.

8 Conclusion

Stochastic processes offer effective tools for modeling and analyzing systems with natural randomness, from discrete Markov chains to continuous Brownian motion. Concepts like martingales and Ito calculus support accurate analysis of uncertainty in fields such as finance, physics, and engineering. These methods are useful for solving complex real-world problems.